



# Efficient programs

## *myJoin*

Mihnea Aleman - 12333410

Ivan Babiy - 12142160

Artur Ohanian - 12239164

Jakob Petermandl - 11776823

Marijana Petojevic - 12017529

Alex Polner - 12024704

Paul Spörker - 11912400

# V0: myJoin Baseline Implementation

- C++ implementation
- Reading files line by line
- 4 nested for-loops
- Multi-stage join using **multimap**

```
multimap<string, string> map1(data1.begin(), data1.end());
multimap<string, string> map2(data2.begin(), data2.end());
multimap<string, string> map3(data3.begin(), data3.end());
multimap<string, string> map4(data4.begin(), data4.end());

for (const auto& [key1, value1] : map1) {
    auto range2 = map2.equal_range(key1);
    auto range3 = map3.equal_range(key1);


    for (auto it2 = range2.first; it2 != range2.second; ++it2) {
        for (auto it3 = range3.first; it3 != range3.second; ++it3) {
            auto range4 = map4.equal_range(it3->second);

            for (auto it4 = range4.first; it4 != range4.second; ++it4) {
                cout << it3->second << "," << key1 << "," << value1 << "," << it2->second << "," << it4->second << "\n";
            }
        }
    }
}
```

Number of cycles: 535.027.663.415

# V1: Compiler flags

- Compiled with different flags
  - -O2
  - -O3
  - -Ofast



Number of cycles -O2: 353,151,865,678

Number of cycles -O3: 355,163,021,185

Number of Cycles -Ofast: 354,211,527,370

## V2: Improved line splitting

- Avoid creating a `stringstream`
- `find` and `substr`
- Use `emplace_back` instead of `push_b`

```
// Split a string into tokens by a given delimiter
vector<string> split(const string &line, char delimiter) {
    vector<string> tokens;
    string token;
    size_t start = 0, end = 0;

    while ((end = line.find(c: delimiter, pos: start)) != string::npos) {
        tokens.emplace_back(line.substr(pos: start, n: end - start));
        start = end + 1;
    }
    tokens.emplace_back(line.substr(pos: start)); // Add the last token
    return tokens;
}
```

Number of cycles: 312 277 715 826

## V3: Use `std::move`

- Use `std::move`
- Transfers resource from one object to another
- Avoids copying the string

```
auto tokens : vector<string> = split(line, delimiter: ',');  
if (tokens.size() == 2) {  
    data.emplace_back(t1: std::move(tokens[0]), t2: std::move(tokens[1]));  
}
```

Number of cycles: 308 833 085 225

## V4: Use unordered\_multimap

- Use `unordered_multimap` instead of `multimap`
- Insert and search in average  $O(1)$  instead of  $O(\log(n))$

```
unordered_multimap<string, string> map1;  
unordered_multimap<string, string> map2;  
unordered_multimap<string, string> map3;  
unordered_multimap<string, string> map4;
```

Number of cycles: 121 187 831 303

# V5: Mem reservations and String scope

- Based on myjoin4
- Reserve memory for `std::vector` in `read_file`
- Reduce scope of `std::string`
- Remove now unused `split` method

```
vector<pair<string, string>> read_file(const string &filename) {  
    vector<pair<string, string>> data;  
    ifstream file(filename);  
    string line;  
  
    file.seekg(0, ios::end); // Estimate file size for pre-allocat  
    size_t estimated_lines = file.tellg() / 50; // Rough guess: 50  
    file.seekg(0, ios::beg);  
  
    data.reserve(estimated_lines); // Reserve space  
  
    while (getline(file, line)) {  
        //reduce scope of std::string  
        size_t delim = line.find(',');  
        if (delim != string::npos) {  
            data.emplace_back(  
                line.substr(0, delim),  
                line.substr(delim + 1)  
            );  
        }  
    }  
}
```

Number of cycles: 109 042 856 561

# V6: Loop unrolling

- Based on myjoin5
- Unroll the two **innermost** loops by **two**
- Reduction of the branch overhead as well as the overall iterations
- Use a buffer for the output and flush it in one go for a more efficient output

```
std::ostringstream buffer; // Use a buffer for efficient output

for (const auto &[key1, value1] : map1) {
    auto range2 = map2.equal_range(key1);
    auto range3 = map3.equal_range(key1);
    if (range2.first == range2.second || range3.first == range3.second) continue; //

    for (auto it2 = range2.first; it2 != range2.second; ++it2) {
        auto it3 = range3.first;
        while (it3 != range3.second) {
            // Unroll by 2 for it3 loop
            if (it3 != range3.second) {
                auto range4 = map4.equal_range(it3->second);
                auto it4 = range4.first;
                while (it4 != range4.second) {
                    // Unroll by 2 for it4 loop
                    if (it4 != range4.second) {
                        buffer << it3->second << "," << key1 << "," << value1 << ","
                            << it2->second << "," << it4->second << "\n";
                        ++it4;
                    }
                    if (it4 != range4.second) {
                        buffer << it3->second << "," << key1 << "," << value1 << ","
                            << it2->second << "," << it4->second << "\n";
                        ++it4;
                    }
                }
                ++it3;
            }
        }
    }
}
```

Number of cycles: 105 060 222 941



# V7: Parallelize csv reading

- Based on myjoin6
- Use **async-threads** to parallelize IO
- Utilizing available resources effectively
- Reducing the backend overhead
- Only decreased the time not cycles

```
void my_join(const string &file1, const string &file2, const string &file3,
// read files in parallel
    auto future_data1 = std::async(std::launch::async, read_file, file1);
    auto future_data2 = std::async(std::launch::async, read_file, file2);
    auto future_data3 = std::async(std::launch::async, read_file, file3);
    auto future_data4 = std::async(std::launch::async, read_file, file4);

    // wait for files to be read
    auto data1 = future_data1.get();
    auto data2 = future_data2.get();
    auto data3 = future_data3.get();
    auto data4 = future_data4.get();

    unordered_multimap<string, string> map1, map2, map3, map4;
```

Number of cycles: 113 064 693 638

# V8: Initialize unordered\_multimap w/ reserved space

- Based on myjoin6
- Removed creation of **map1** - map not necessary for file1
- Initialized **unordered\_multimap** with reserved space

```
void my_join(const string &file1, const string &file2, const string &file3, co
    auto data1 = read_file(file1);
    auto data2 = read_file(file2);
    auto data3 = read_file(file3);
    auto data4 = read_file(file4);

    // Initialize unordered_multimap with reserved space to reduce rehashing
    unordered_multimap<string, string> map2(data2.begin(), data2.end());
    unordered_multimap<string, string> map3(data3.begin(), data3.end());
    unordered_multimap<string, string> map4(data4.begin(), data4.end());
```

Number of cycles: 90 858 019 300

# V9: Parallelize join process

- Use `std::thread`
- Process different chunks of `data1` in parallel
- Merge the results into output buffer
- Cycles increased, time decreased

Number of cycles: 99 157 079 419

```
void my_join(const string &file1, const string &file2, const string &file3, const string &file4) {
    auto data1 = read_file(file1);
    auto data2 = read_file(file2);
    auto data3 = read_file(file3);
    auto data4 = read_file(file4);

    // Initialize unordered_multimap with reserved space to reduce rehashing
    unordered_multimap<string, string> map2(data2.begin(), data2.end());
    unordered_multimap<string, string> map3(data3.begin(), data3.end());
    unordered_multimap<string, string> map4(data4.begin(), data4.end());

    // prepare multithreading
    size_t num_threads = thread::hardware_concurrency();
    if (num_threads == 0) num_threads = 4; // default 4 threads
    size_t chunk_size = data1.size() / num_threads;

    vector<thread> threads;
    vector<vector<pair<string, string>>> chunks;

    // split data1 into chunks
    for (size_t i = 0; i < num_threads; ++i) {
        auto start = data1.begin() + i * chunk_size;
        auto end = (i == num_threads - 1) ? data1.end() : start + chunk_size;
        chunks.emplace_back(start, end);
    }

    ostringstream output_buffer;
    mutex output_mutex;

    // launch threads
    for (size_t i = 0; i < num_threads; ++i) {
        threads.emplace_back(process_chunk, ref(chunks[i]), ref(map2), ref(map3), ref(map4),
                             ref(output_buffer), ref(output_mutex));
    }

    // wait for threads to complete
    for (auto &t : threads) {
        t.join();
    }

    // Flush the buffer to standard output
    cout << output_buffer.str();
}
```

# V10: Better output handling

- Based on myjoin8
- Write output in batches - instead of storing results in memory
- Flush buffer to `std::cout` when threshold is reached
- Reduces memory usage
- Cycles increased, time unchanged

```
void my_join(const string &file1, const string &file2, const string &file3,
             for (const auto &[key1, value1] : data1) {
    // check matches in map2 and map3
    auto range2 = map2.equal_range(key1);
    auto range3 = map3.equal_range(key1);

    if (range2.first == range2.second || range3.first == range3.second)
        continue;

    for (auto it2 = range2.first; it2 != range2.second; ++it2) {
        for (auto it3 = range3.first; it3 != range3.second; ++it3) {
            // lookup map4 matches for current value from map3
            auto range4 = map4.equal_range(it3->second);
            for (auto it4 = range4.first; it4 != range4.second; ++it4)
                output_buffer << it3->second << "," << key1 << "," << value1 << "\n";
            output_buffer << it2->second << "," << it4->second << "\n";

            // track buffer size
            current_buffer_size += output_buffer.str().size();

            // flush buffer to stdout if it exceeds threshold
            if (current_buffer_size >= BUFFER_THRESHOLD) {
                cout << output_buffer.str();
                output_buffer.str(""); // clear buffer
                output_buffer.clear(); // reset error state
                current_buffer_size = 0; // reset size counter
            }
        }
    }
}

// Flush any remaining content in the buffer
if (current_buffer_size > 0) {
    cout << output_buffer.str();
}
```

Number of cycles: 95 875 908 602

# V11: In memory sorting

- Based on myjoin8
- Sort **data2**, **data3**, and **data4** by their keys
- Allows binary search instead of hash based lookups
- Reduces random memory access
- Cycles and time increased (a lot)

Number of cycles: 280 943 563 629

```
void my_join(const string &file1, const string &file2, const string &file3, const
auto data1 = read_file(file1);
auto data2 = read_file(file2);
auto data3 = read_file(file3);
auto data4 = read_file(file4);

// sort data2, data3, data4 by key
sort(data2.begin(), data2.end(), [](const auto &a, const auto &b) {
    return a.first < b.first;
});

sort(data3.begin(), data3.end(), [](const auto &a, const auto &b) {
    return a.first < b.first;
});

sort(data4.begin(), data4.end(), [](const auto &a, const auto &b) {
    return a.first < b.first;
});

ostringstream output_buffer;

for (const auto &[key1, value1] : data1) {
    // binary search for matches in data2 and data3
    auto matches2 = find_all_matches(data2, key1);
    auto matches3 = find_all_matches(data3, key1);

    if (matches2.empty() || matches3.empty())
        continue;

    for (const auto &[key2, value2] : matches2) {
        for (const auto &[key3, value3] : matches3) {
            // binary search for matches in data4
            auto matches4 = find_all_matches(data4, value3);

            for (const auto &[key4, value4] : matches4) {
                output_buffer << value3 << "," << key1 << "," << value1 << ","
                    << value2 << "," << value4 << "\n";
            }
        }
    }
}

// Flush the buffer to standard output
cout << output_buffer.str();
}
```

# V12: Minimize string operations

- Based on myjoin8
- Assign unique **int** IDs to strings (i.e.: “intern” strings)
- Lower memory overhead during joins
- Cycles and time increased

```
// interning helper: maps strings to unique int IDs
class StringInterner {
private:
    unordered_map<string, int> string_to_id;
    vector<string> id_to_string;
    int next_id = 0;

public:
    // get or assign ID to string
    int intern(const string &s) {
        auto it = string_to_id.find(s);
        if (it != string_to_id.end()) {
            return it->second;
        }
        string_to_id[s] = next_id;
        id_to_string.push_back(s);
        return next_id++;
    }

    // get original string from ID
    const string &resolve(int id) const {
        return id_to_string[id];
    }
};
```

Number of cycles: 120 568 790 850

# V13: Parallelize output writing

- Based on myjoin8
- Offload output writing to separate thread
- Decouples output writing (I/O) from join process
- Cycles and time increased

Number of cycles: 123 938 714 878

```
void my_join(const string &file1, const string &file2, const string &file3, const string &
// start writer thread
thread writer(writer_thread_func);

for (const auto &[key1, value1] : data1) {
    auto range2 = map2.equal_range(key1);
    auto range3 = map3.equal_range(key1);

    if (range2.first == range2.second || range3.first == range3.second)
        continue;

    for (auto it2 = range2.first; it2 != range2.second; ++it2) {
        for (auto it3 = range3.first; it3 != range3.second; ++it3) {
            auto range4 = map4.equal_range(it3->second);
            for (auto it4 = range4.first; it4 != range4.second; ++it4) {
                // enqueue the output line
                {
                    lock_guard<mutex> lock(queue_mutex);
                    output_queue.push(it3->second + "," + key1 + "," + value1 + "," +
                                     it2->second + "," + it4->second + "\n");
                }
                cv.notify_one();
            }
        }
    }
}

// signal writer thread to finish
{
    lock_guard<mutex> lock(queue_mutex);
    finished = true;
}
cv.notify_one();
writer.join();
}
```

# V14: Use memory-mapped files

- Based on myjoin8
- Used **mmap** for reading **.csv** files
- Reduces overhead of file I/O
- OS handles paging more efficiently
- Cycles and time remained the same

```
vector<pair<string, string>> read_file(const string &filename) {  
    size_t start = 0;  
    for (size_t i = 0; i < file_size; ++i) {  
        if (file_in_memory[i] == '\n') {  
            size_t line_length = i - start;  
            string line(file_in_memory + start, line_length);  
            size_t delim = line.find(',');  
            if (delim != string::npos) {  
                data.emplace_back(  
                    line.substr(0, delim),  
                    line.substr(delim + 1)  
                );  
            }  
            start = i + 1;  
        }  
    }  
  
    // handle last line (doesn't end with newline)  
    if (start < file_size) {  
        string line(file_in_memory + start, file_size - start);  
        size_t delim = line.find(',');  
        if (delim != string::npos) {  
            data.emplace_back(  
                line.substr(0, delim),  
                line.substr(delim + 1)  
            );  
        }  
    }  
  
    munmap(file_in_memory, file_size);  
    close(fd);  
    return data;  
}
```

Number of cycles: 92 337 135 105



# V15: Further optimize data structure

- Based on myjoin8
- Replace `unordered_multimap` with `robin_hood::unordered_map`
- `robin_hood` hashing offer faster hash maps with reduced memory overhead

```
void my_join(const string &file1, const string &file2, const
robin_hood::unordered_map<string, vector<string>> map2;
robin_hood::unordered_map<string, vector<string>> map3;
robin_hood::unordered_map<string, vector<string>> map4;

map2.reserve(data2_size);
map3.reserve(data3_size);
map4.reserve(data4_size);

// Populate maps
for (const auto& [key, value] : data2) {
    map2[key].push_back(value);
}
for (const auto& [key, value] : data3) {
    map3[key].push_back(value);
}
for (const auto& [key, value] : data4) {
    map4[key].push_back(value);
}
```

Number of cycles: 85 025 197 043

# V16: String view

- Based on myjoin15
- Use `std::string_view` instead of `std::string` - to avoid unnecessary memory allocations
- Read the entire file into memory at once - instead of reading line by line - to reduce I/O operations

```
vector<pair<string_view, string_view>> read_file(const string &filename, string
ifstream file(filename, ios::ate | ios::binary);
if (!file) {
    cerr << "Error: Unable to open file " << filename << "\n";
    exit(1);
}

// Read the entire file into memory
size_t file_size = file.tellg();
file.seekg(0, ios::beg);
file_content.resize(file_size);
file.read(file_content.data(), file_size);

// Parse the file content into key-value pairs
vector<pair<string_view, string_view>> data;
size_t estimated_lines = file_size / 50; // estimate
data.reserve(estimated_lines);

size_t start = 0;
size_t end = 0;
while ((end = file_content.find('\n', start)) != string::npos) {
    size_t delim = file_content.find(',', start);
    if (delim != string::npos && delim < end) {
        data.emplace_back(
            string_view(file_content.data() + start, delim - start),
            string_view(file_content.data() + delim + 1, end - delim - 1));
    }
    start = end + 1;
}

return data;
}
```

Number of cycles: 46 191 998 137

# V17:

- Based on myjoin16
- Removed `ostreamstream`, it can be less efficient than `string`
- Output in chunks of 64KB

```
if (output_buffer.size() >= 65536) {  
    cout << output_buffer;  
    output_buffer.clear();  
    output_buffer.reserve(65536);  
}
```

Number of cycles: 40 973 671 158

# V18: V17 with parallel file reading

- Parallel file reading with `std::launch`
- Decreased time but increased time number of iterations

```
void my_join(const string& file1, const string& file2,  
            const string& file3, const string& file4) {  
    // Parallel file reading  
    auto future1 = async(launch::async, read_file, file1);  
    auto future2 = async(launch::async, read_file, file2);  
    auto future3 = async(launch::async, read_file, file3);  
    auto future4 = async(launch::async, read_file, file4);  
  
    auto f1 = future1.get();  
    auto f2 = future2.get();  
    auto f3 = future3.get();  
    auto f4 = future4.get();  
}
```


Number of cycles: 47 498 081 175

# V18 vs. V17

	V17	V18 (3 CPUs)
Task-clock (s)	11.22	13.98
Time-elapsed (s)	11.22	4.65
Cycles (B)	40.97	47.48


# Results

Benchmark Results - Cycles




# Results

Benchmark Results - Time



# Single Optimizations







# Thanks for your attention!

Questions?